



# **Project | Deep Learning: Image Classification using CNN and Transfer Learning**

Presented by:

Mohamad Traiki  
Sergei Volkov  
Mitesh Parab

# Dataset Overview and Objective

**Dataset used:** CIFAR Dataset (CIFAR-10 / CIFAR-100)

- Created by **Alex Krizhevsky, Vinod Nair, Geoffrey Hinton**
  - **60,000 colored images (32×32) across 10 classes** (CIFAR-10)
  - Classes include: airplane, automobile, bird, cat, deer, dog, frog, horse, ship, truck
- 
- To develop and evaluate image classification models using **Convolutional Neural Networks (CNNs)** and **Transfer Learning** on the **CIFAR dataset**.
  - To compare custom-built CNN architectures with pre-trained deep learning models, analyzing their performance, accuracy, and efficiency in classifying images.

Used **pre-trained models** from Keras Applications (e.g., EfficientNet B0 to B7, MobileNetV2 etc..)

Used Optimizers (e.g., Adam, RMSprop, Adadelta etc..)

# Data pre-processing

## Self-made CNNs

Load data (cifar10)  Change to grayscale  Normalize

 Convert labels to one-hot vectors

## Transfer Models

Load data (cifar10)  Resize from 32 x 32 to 96 x 96



# Callbacks

## Early Stopping:

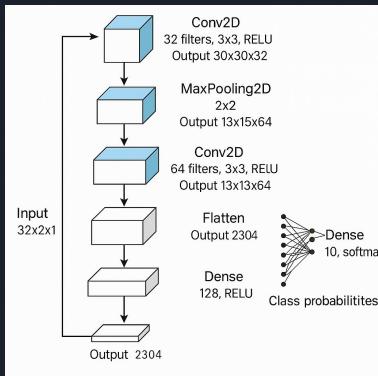
- Monitors validation accuracy
- Patience of 8 epochs. Terminates training when no improvement noticed
- Saves parameters from best performing epoch
- Early Stopping was performed on all our Models.

## Learning Rate Reduction: (Transfer Models)

- Monitors validation accuracy
- Patience of 4 epochs
- Reduces learning rate by factor  $f$  when no improvement noticed.
- Sets a minimum learning rate that can't be reduced.
- LRR was performed on only Transfer Models

# Self-made CNNs

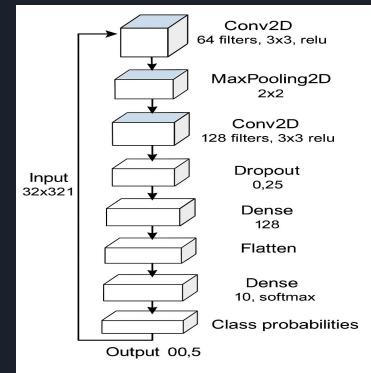
Initial Model (Adam):



- Same number of layers.
- CNN with RMS: double number of filters
- CNN with RMS: dropout 25% of neurons.
- Both CNNS: average performance

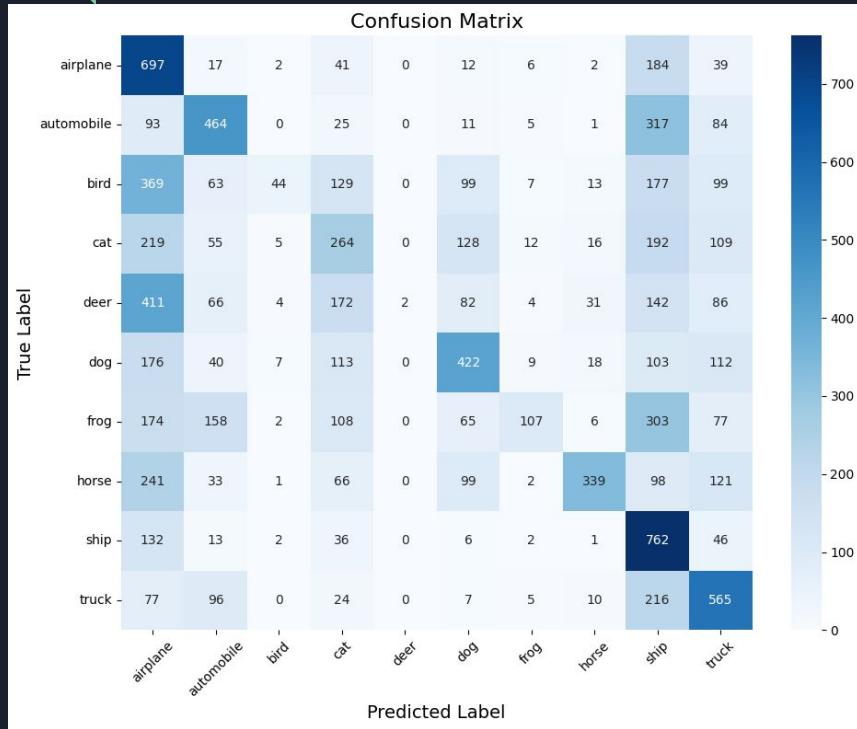
Activation: relu  
Optimizer: Adam

CNN with RMSprop

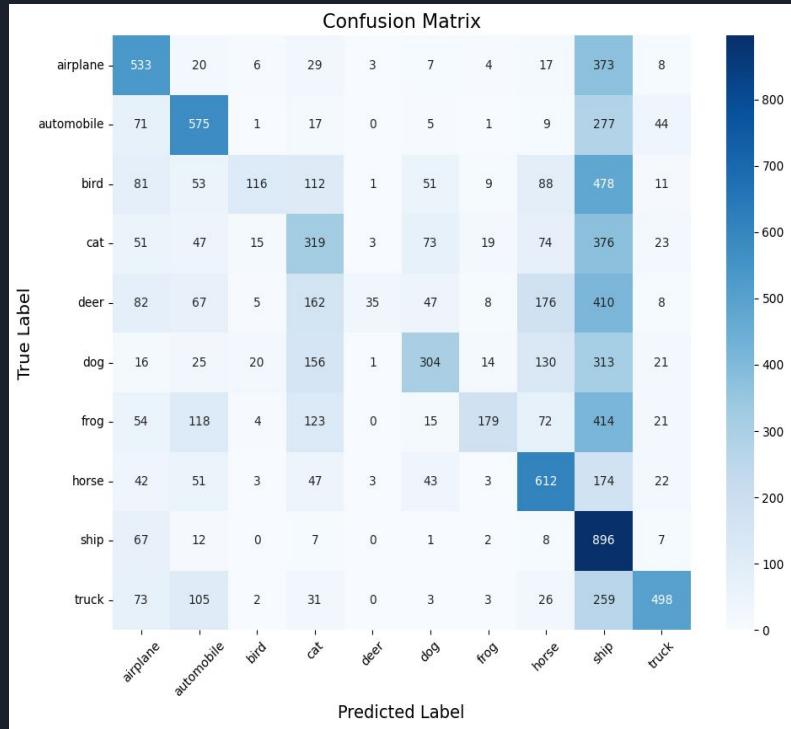


Activation: relu  
Optimizer: RMSprop

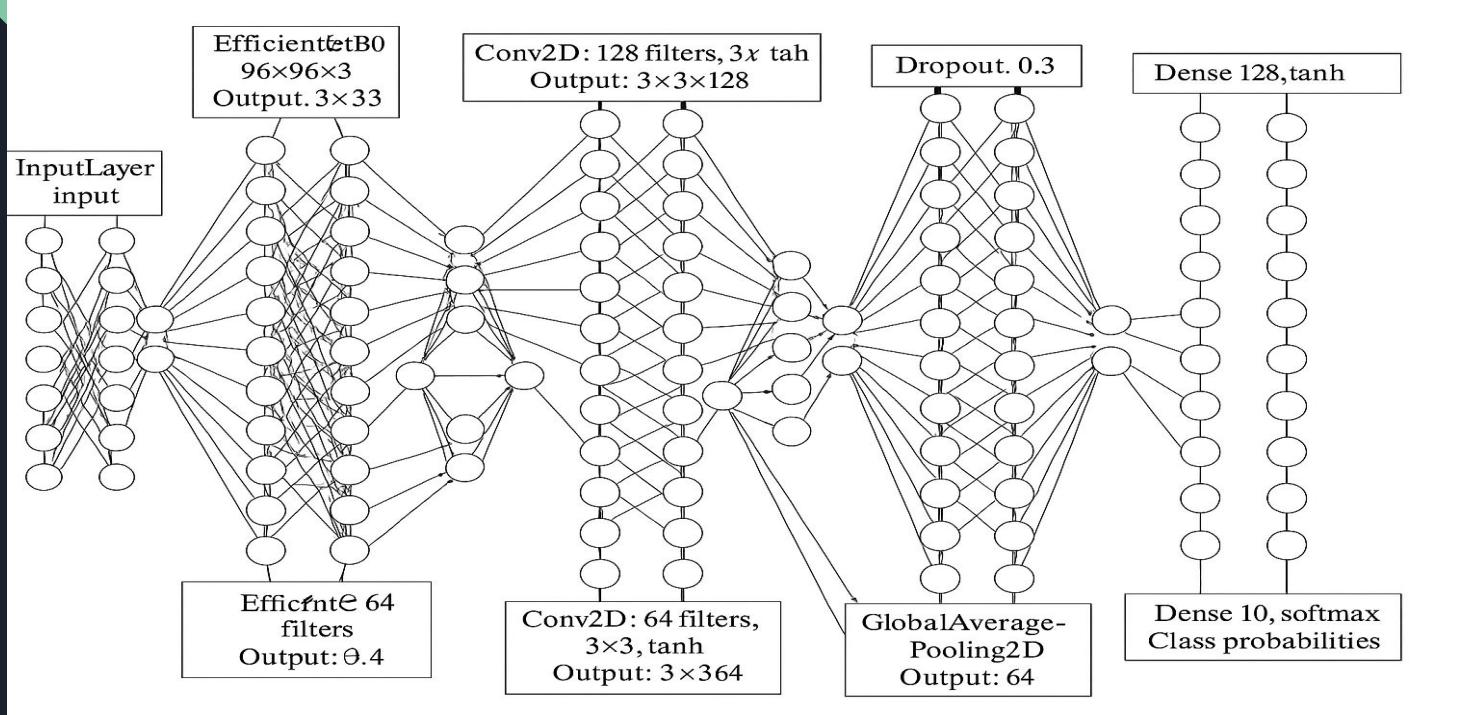
## Initial Model (Adam)



## CNN (RMSprop)



# Transfer Models

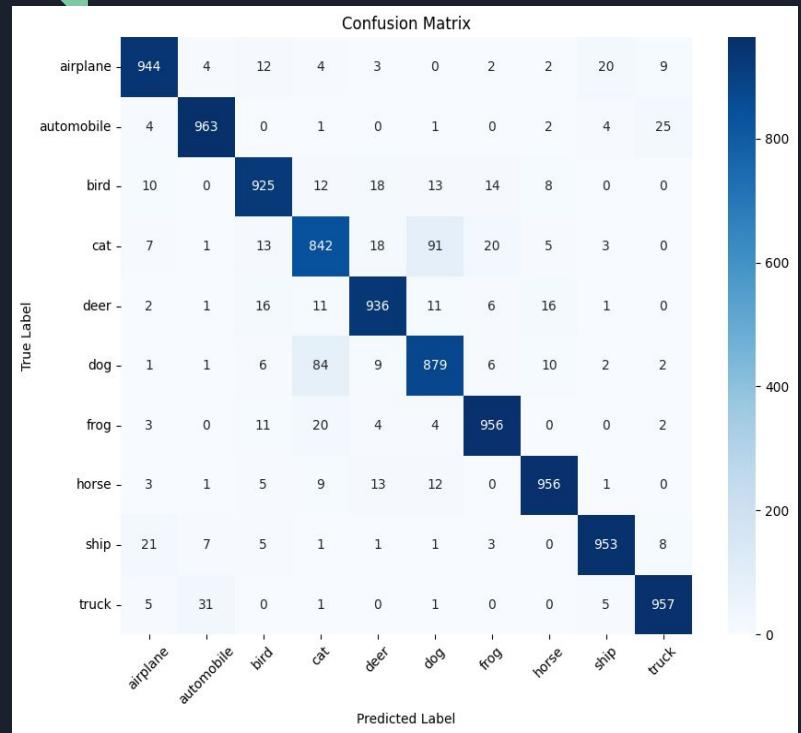




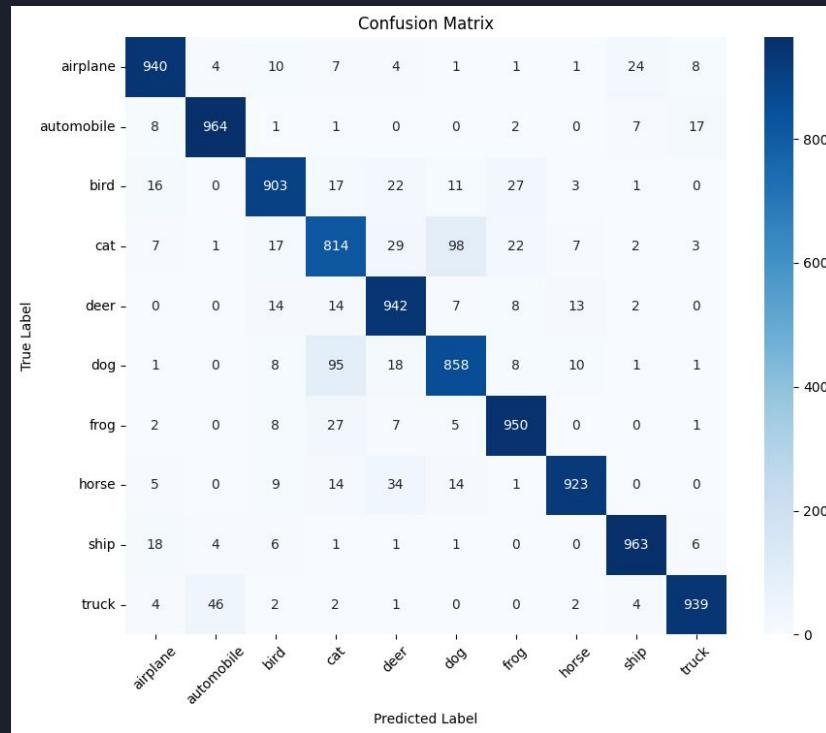
## Transfer Models Architecture:

- Base model + Conv2D(128) + dropout(0.3) + Conv2D(64) + Global AVG Pooling + Dense(128) + dropout (0.4) + Dense(10)(Output Layer with 10 Classifiers)
- Both models share same architecture
- Differ in activation function (relu vs tanh)
- Use Early Stopping and Learn Rate Reduction
- Early stopping:
  - Patience: 8 epochs
  - Monitor: val\_accuracy
  - Saves best weights
- Learning Rate Reduction:
  - Patience: 4 epochs
  - Monitor: val\_accuracy
  - Factor: 0.2
- Introduce Checkpoints
- Initial Learning Rate: 0.001
- Fine tuned by freezing the initial layers of base model and training last 50 layers.

EfficientNETB0(tanh): Test-accuracy 0.920



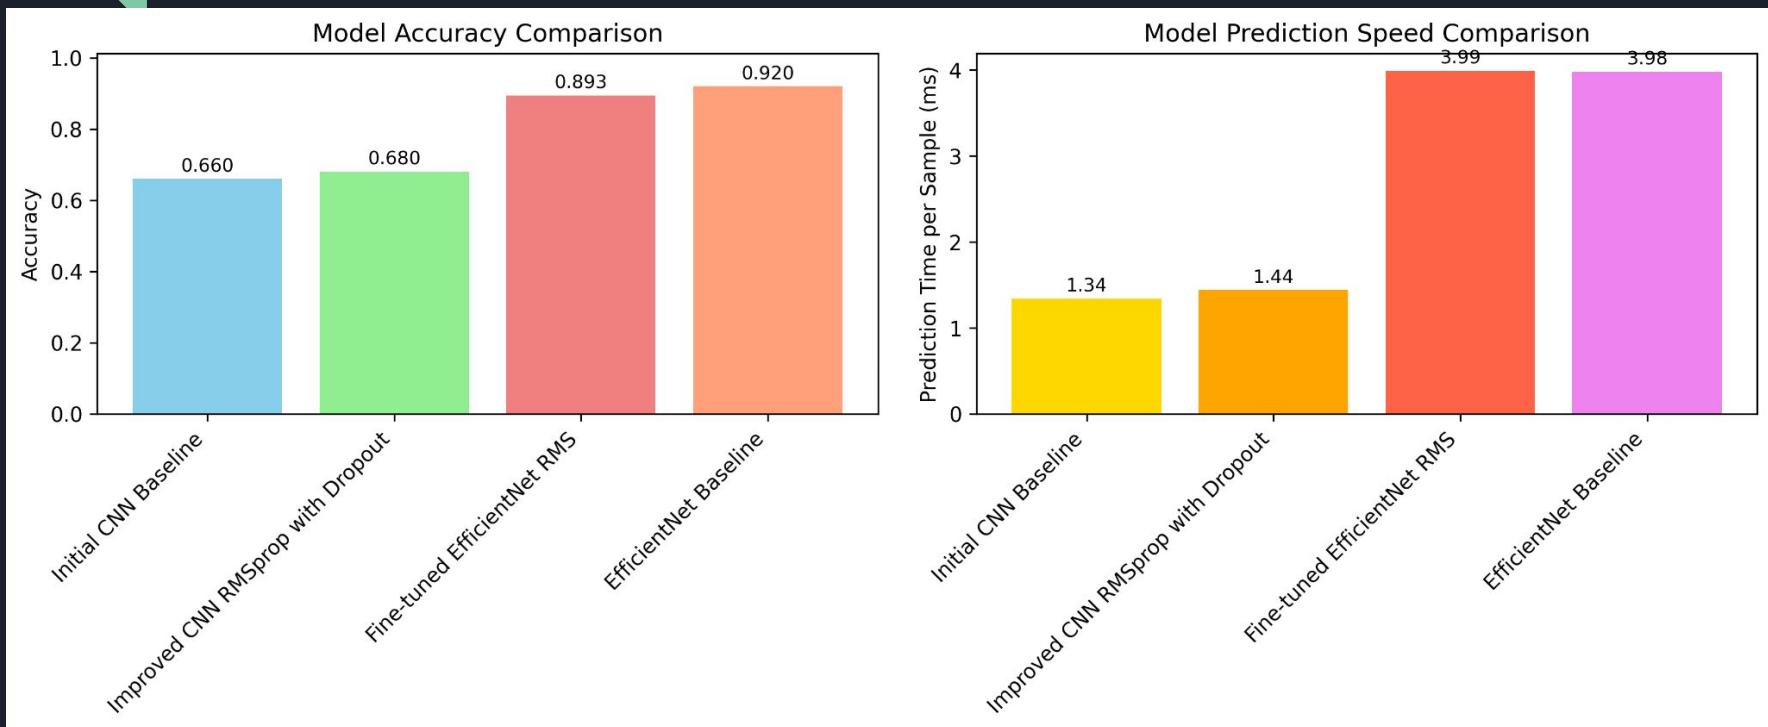
EfficientNETB0(relu): Test-accuracy 0.916





# Models serialization issues

- Be careful saving and loading models between environments
  - Colab uses Keras 3.10 which sometimes isn't compatible with latest 3.12
  - This may lead you to corrupted/unreadable model file
- Don't use Lambda layers for preprocessing
  - Some preprocessing layers may not be loaded from .keras files (probably not limited to Lambda)
  - This will result to unreadable model even in the same env
- Reproducibility in Python ML libraries is complicated



# Performance comparison

Model Name	Accuracy	Precision	Recall	F1-Score	Loss	Pred Time/ms
-----						
Initial CNN Baseline	0.6599	0.6600	0.6599	0.6590	1.0501	1.3399
Improved CNN RMSprop with Dropout	0.6799	0.6865	0.6799	0.6781	1.0629	1.4427
Fine-tuned EfficientNet RMS	0.8935	0.8935	0.8935	0.8935	0.3988	3.9895
EfficientNet Baseline	0.9196	0.9199	0.9196	0.9196	0.6340	3.9787



T - Transfer Learning

H - Hyperparameters

R - ReLu Activation

A - Augmentation

N - Neurons

K - Kernel (Convolutional)

YOU!